

浙江大学实验报告

课程名称： 操作系统 实验类型： 综合型
实验项目名称： 添加系统调用
学生姓名： 沈子衿 学号： 3160104734
电子邮件地址： zijinshen@zju.edu.cn
实验日期： 2018 年 12 月 22 日

一、实验环境

主机配置：

- 主机型号：Lenovo ThinkPad T480
- 内存：16GB
- 处理器：i7-8550U
- 操作系统：Windows 10 家庭中文版

虚拟机配置：

- 虚拟机环境：Vmware Workstation Pro 14
- Ubuntu 版本：16.04
- Linux 内核版本：4.8

二、实验内容和结果及分析


1. 实验设计思路

本次实验的设计思路已经在实验参考中给出，总体还是较为清晰的，思路如下：

- 1) 下载并解压内核；
- 2) 为内核打补丁；
- 3) 配置内核；
- 4) 添加新的系统调用号（本次实验我们使用 223）；
- 5) 在系统调用表中添加或修改响应表项，以便系统调用处理程序检索；
- 6) 修改统计系统缺页次数和进程缺页次数的内核代码；
- 7) 实现简单的系统调用（`sys_mysyscall`）；
- 8) 编译内核和重启内核；
- 9) 撰写对应的用户态程序，输出结果。

2. 实验步骤及截图

首先，从阿里云镜像下载 linux-4.8 版本内核和对应补丁，存放在主目录（~）下：



输入命令 `tar -xvf linux-4.8.tar.xz` 解压内核:

```
danielshen@ubuntu: ~  
linux-4.8/virt/kvm/arm/vgic/vgic-its.c  
linux-4.8/virt/kvm/arm/vgic/vgic-kvm-device.c  
linux-4.8/virt/kvm/arm/vgic/vgic-mmio-v2.c  
linux-4.8/virt/kvm/arm/vgic/vgic-mmio-v3.c  
linux-4.8/virt/kvm/arm/vgic/vgic-mmio.c  
linux-4.8/virt/kvm/arm/vgic/vgic-mmio.h  
linux-4.8/virt/kvm/arm/vgic/vgic-v2.c  
linux-4.8/virt/kvm/arm/vgic/vgic-v3.c  
linux-4.8/virt/kvm/arm/vgic/vgic.c  
linux-4.8/virt/kvm/arm/vgic/vgic.h  
linux-4.8/virt/kvm/async_pf.c  
linux-4.8/virt/kvm/async_pf.h  
linux-4.8/virt/kvm/coalesced_mmio.c  
linux-4.8/virt/kvm/coalesced_mmio.h  
linux-4.8/virt/kvm/eventfd.c  
linux-4.8/virt/kvm/irqchip.c  
linux-4.8/virt/kvm/kvm_main.c  
linux-4.8/virt/kvm/vfio.c  
linux-4.8/virt/kvm/vfio.h  
linux-4.8/virt/lib/  
linux-4.8/virt/lib/Kconfig  
linux-4.8/virt/lib/Makefile  
linux-4.8/virt/lib/irqbypass.c  
danielshen@ubuntu:~$
```

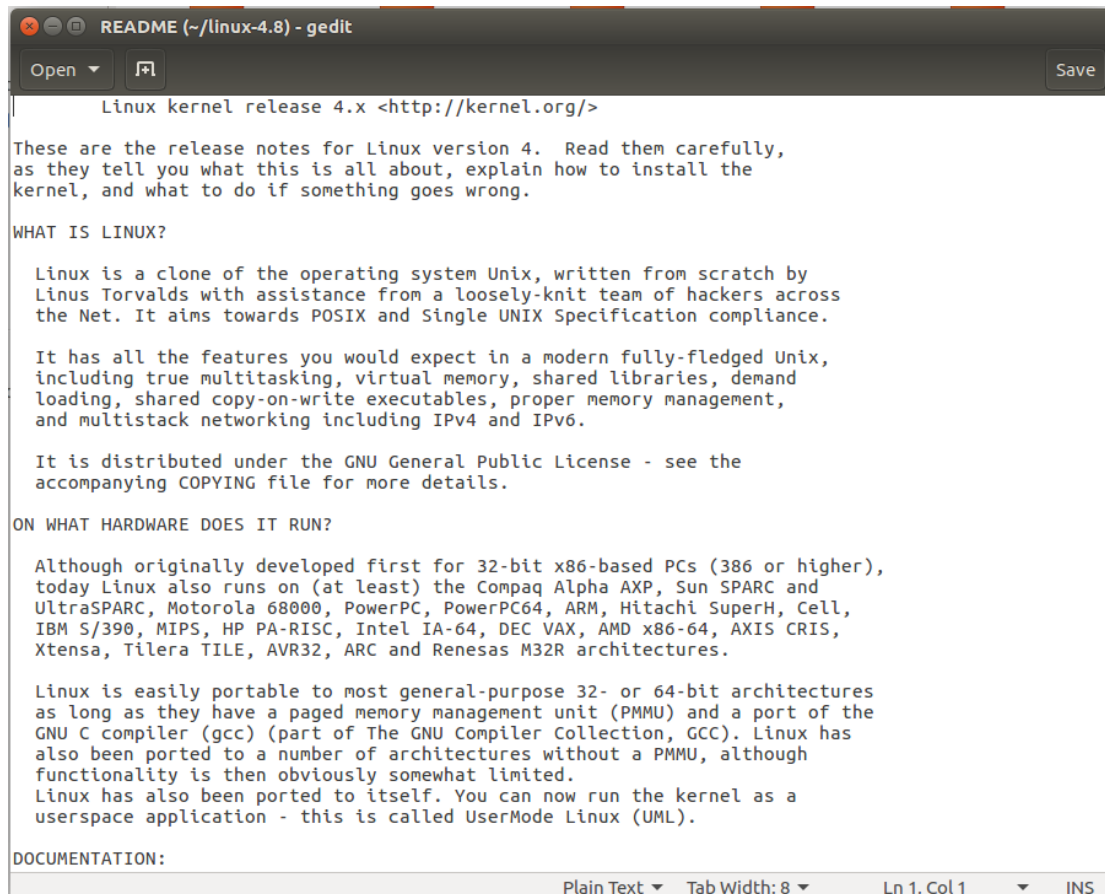
输入命令 `xz -d patch-4.8.xz | patch -p1` 为内核打补丁

```
linux-4.8/virt/lib/Kconfig  
linux-4.8/virt/lib/Makefile  
linux-4.8/virt/lib/irqbypass.c  
danielshen@ubuntu:~$ xz -d patch-4.8.xz | patch -p1  
danielshen@ubuntu:~$
```

在 ubuntu 环境下, 用命令 `make menuconfig` 对内核进行配置时, 需要用终端模式下的字符菜单支持软件包 `libncurses5-dev`, 因此你是第一次重建内核, 需要下载并安装该软件包, 下载并安装命令如下: `apt-get install libncurses5-dev libssl-dev`:

```
danielshen@ubuntu:~$ sudo apt-get install libncurses5-dev libssl-dev  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
The following additional packages will be installed:  
  libssl-doc libtinfo-dev zlib1g-dev  
Suggested packages:  
  ncurses-doc  
The following NEW packages will be installed:  
  libncurses5-dev libssl-dev libssl-doc libtinfo-dev zlib1g-dev  
0 upgraded, 5 newly installed, 0 to remove and 92 not upgraded.  
Need to get 2,843 kB of archives.  
After this operation, 12.0 MB of additional disk space will be used.  
Do you want to continue? [Y/n]
```

安装完成后, 查看内核 README 文件:



进入 linux-4.8 目录，输入 make mrproper 命令清空缓存：

```
danielshen@ubuntu:~$ cd ~/.linux-4.8
danielshen@ubuntu:~/linux-4.8$ make mrproper
danielshen@ubuntu:~/linux-4.8$
```

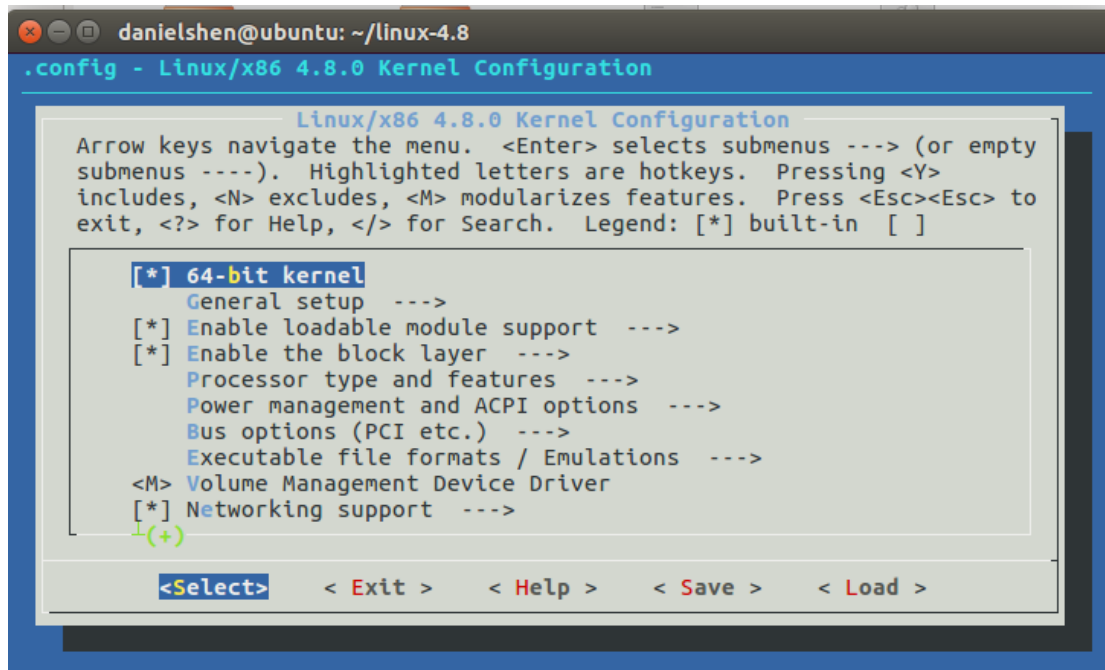
输入 cp /boot/config-`uname -r` .config 命令复制配置文件到内核根目录，将新内核配置与正在运行的操作系统内核的运行环境匹配，然后 make menuconfig：

```
danielshen@ubuntu:~/linux-4.8$ make menuconfig
HOSTCC scripts/basic/fixdep
HOSTCC scripts/kconfig/mconf.o
SHIPPED scripts/kconfig/zconf.tab.c
SHIPPED scripts/kconfig/zconf.lex.c
SHIPPED scripts/kconfig/zconf.hash.c
HOSTCC scripts/kconfig/zconf.tab.o
HOSTCC scripts/kconfig/lxdialog/checklist.o
HOSTCC scripts/kconfig/lxdialog/util.o
HOSTCC scripts/kconfig/lxdialog/inputbox.o
HOSTCC scripts/kconfig/lxdialog/textbox.o
HOSTCC scripts/kconfig/lxdialog/yesno.o
HOSTCC scripts/kconfig/lxdialog/menubox.o
HOSTLD scripts/kconfig/mconf
scripts/kconfig/mconf Kconfig
.config:4237:warning: symbol value 'm' invalid for GPIO_MB86S7X
configuration written to .config

*** End of the configuration.
*** Execute 'make' to start the build or try 'make help'.

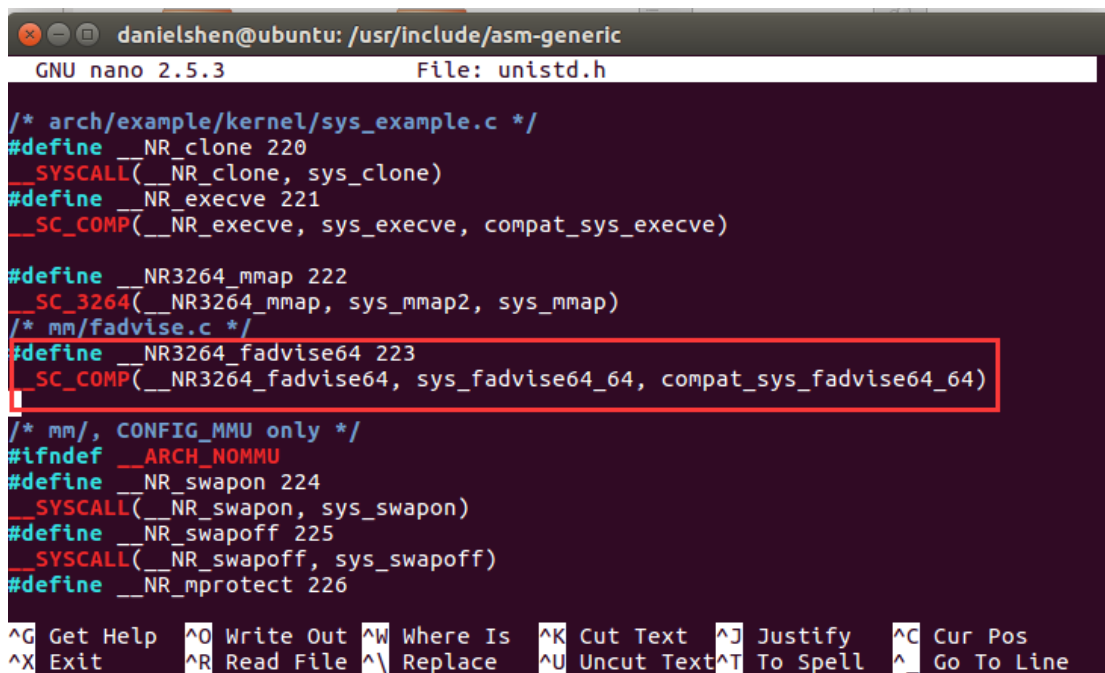
danielshen@ubuntu:~/linux-4.8$
```

期间命令行会弹出一个配置界面，这里直接选择“exit”并“save”即可：



这之后，分别在本系统的 unistd.h 和新内核的 unistd.h 中添加系统调用号：

修改前：



修改号：

```
danielshen@ubuntu: /usr/include/asm-generic
GNU nano 2.5.3 File: unistd.h Modified

/* arch/example/kernel/sys_example.c */
#define __NR_clone 220
__SYSCALL(__NR_clone, sys_clone)
#define __NR_execve 221
__SC_COMP(__NR_execve, sys_execve, compat_sys_execve)

#define __NR3264_mmap 222
__SC_3264(__NR3264_mmap, sys_mmap2, sys_mmap)
/* mm/fadvise.c */
#define __NR_mysyscall 223
__SYSCALL(__NR_mysyscall, sys_mysyscall)

/* mm/, CONFIG_MMU only */
#ifndef __ARCH_NOMMU
#define __NR_swapon 224
__SYSCALL(__NR_swapon, sys_swapon)
#define __NR_swapoff 225
__SYSCALL(__NR_swapoff, sys_swapoff)
#define __NR_mprotect 226
```

对新内核中的 unistd.h 也作同样修改：

```
danielshen@ubuntu:~/linux-4.8$ cd ./include/uapi/asm-generic/
danielshen@ubuntu:~/linux-4.8/include/uapi/asm-generic$
```

```
danielshen@ubuntu: ~/linux-4.8/include/uapi/asm-generic
GNU nano 2.5.3 File: unistd.h

#define __NR_keyctl 219
__SC_COMP(__NR_keyctl, sys_keyctl, compat_sys_keyctl)

/* arch/example/kernel/sys_example.c */
#define __NR_clone 220
__SYSCALL(__NR_clone, sys_clone)
#define __NR_execve 221
__SC_COMP(__NR_execve, sys_execve, compat_sys_execve)

#define __NR3264_mmap 222
__SC_3264(__NR3264_mmap, sys_mmap2, sys_mmap)
/* mm/fadvise.c */
#define __NR_mysyscall 223
__SYSCALL(__NR_mysyscall, sys_mysyscall)

/* mm/, CONFIG_MMU only */
#ifndef __ARCH_NOMMU
#define __NR_swapon 224
__SYSCALL(__NR_swapon, sys_swapon)

^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos
^X Exit ^R Read File ^\ Replace ^U Uncut Text ^T To Spell ^_ Go To Line
```

上述步骤完成后，修改系统调用表以便系统调用处理程序寻找相应表项：


```

danielshen@ubuntu: ~/linux-4.8/arch/x86/entry/syscalls
danielshen@ubuntu: /usr/include/asm-generic$ nano unistd.h
danielshen@ubuntu: /usr/include/asm-generic$ sudo nano unistd.h
[sudo] password for danielshen:
danielshen@ubuntu: /usr/include/asm-generic$ sudo nano unistd.h
danielshen@ubuntu: /usr/include/asm-generic$ cd ~
danielshen@ubuntu: ~$ ls
Desktop  Downloads  linux-4.8  Music  Pictures  Templates
Documents  examples.desktop  linux-4.8.tar.xz  patch-4.8  Public  Videos
danielshen@ubuntu: ~$ cd linux-4.8
danielshen@ubuntu: ~/linux-4.8$ cd ./include/uapi/asm-generic/
danielshen@ubuntu: ~/linux-4.8/include/uapi/asm-generic$ nano unistd.h
danielshen@ubuntu: ~/linux-4.8/include/uapi/asm-generic$ nano unistd.h
danielshen@ubuntu: ~/linux-4.8/include/uapi/asm-generic$ cd ~
danielshen@ubuntu: ~$ cd linux-4.8/
danielshen@ubuntu: ~/linux-4.8$ ls
arch      crypto      include  kernel      net          security
block     Documentation  init     lib          README       sound
certs     drivers      ipc      lib          REPORTING-BUGS  tools
COPYING   firmware     Kbuild   Makefile     samples      usr
CREDITS   fs           Kconfig  mm           scripts      virt
danielshen@ubuntu: ~/linux-4.8$ cd arch/x86/entry/syscalls/syscall_64.tbl
bash: cd: arch/x86/entry/syscalls/syscall_64.tbl: Not a directory
danielshen@ubuntu: ~/linux-4.8$ cd arch/x86/entry/syscalls/
danielshen@ubuntu: ~/linux-4.8/arch/x86/entry/syscalls$ nano syscall_64.tbl

```

```

GNU nano 2.5.3      File: syscall_64.tbl      Modified:
212  common  lookup_dcookie      sys_lookup_dcookie
213  common  epoll_create        sys_epoll_create
214  64      epoll_ctl_old
215  64      epoll_wait_old
216  common  remap_file_pages    sys_remap_file_pages
217  common  getdents64          sys_getdents64
218  common  set_tid_address     sys_set_tid_address
219  common  restart_syscall     sys_restart_syscall
220  common  semtimedop          sys_semtimedop
221  common  fadvise64           sys_fadvise64
222  64      timer_create        sys_timer_create
223  common  msyscall            sys_msyscall
224  common  timer_gettime       sys_timer_gettime
225  common  timer_getoverrun    sys_timer_getoverrun
226  common  timer_delete        sys_timer_delete
227  common  clock_settime       sys_clock_settime
228  common  clock_gettime       sys_clock_gettime
229  common  clock_getres        sys_clock_getres
230  common  clock_nanosleep     sys_clock_nanosleep

```

此后，在 `include/linux/mm.h` 中声明 `pfcount` 记录总缺页次数：

```
danielshen@ubuntu: ~/linux-4.8/include/linux
GNU nano 2.5.3 File: mm.h Modified

struct user_struct;
struct writeback_control;
struct bdi_writeback;

#ifdef CONFIG_NEED_MULTIPLE_NODES /* Don't use mapnrs, do it properly */
extern unsigned long max_mapnr;
/*modified by zijin*/
extern unsigned long pfcount;
static inline void set_max_mapnr(unsigned long limit)
{
    max_mapnr = limit;
}
#else
static inline void set_max_mapnr(unsigned long limit) { }
#endif

extern unsigned long totalram_pages;
extern void * high_memory;
extern int page_cluster;

[ Unknown Command ]
^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos
^X Exit ^R Read File ^\ Replace ^U Uncut Text ^T To Spell ^_ Go To Line
```

在进程 task_struct 中增加成员 pf，在 include/linux/sched.h 文件中的 task_struct 结构中添加 pf 字段以记录每个进程缺页的次数：

```
danielshen@ubuntu: ~/linux-4.8/include/linux
GNU nano 2.5.3 File: sched.h Modified

    * flushed before IO is initiated or a stale TLB entry potentially
    * allows an update without redirtying the page.
    */
    bool writable;
};

struct task_struct {
    /*modified by zijin*/
    unsigned long pf;
    volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */
    void *stack;
    atomic_t usage;
    unsigned int flags; /* per process flags, defined below */
    unsigned int ptrace;

#ifdef CONFIG_SMP
    struct llist_node wake_entry;
    int on_cpu;
    unsigned int wakee flips;
#endif
}

[XOFF ignored, mumble mumble]
^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos
^X Exit ^R Read File ^\ Replace ^U Uncut Text ^T To Spell ^_ Go To Line
```

修改 dup_task_struct() 函数将子进程的 pf 在复制之后置为 0。注意这一行代码需要放在比较靠后的位置：

```
fork.c (~/.linux-4.8/kernel) - gedit
Open Save

    unsigned long *stackend;

    stackend = end_of_stack(tsk);
    *stackend = STACK_END_MAGIC; /* for overflow detection */
}

static struct task_struct *dup_task_struct(struct task_struct *orig, int node)
{
    struct task_struct *tsk;
    unsigned long *stack;
    int err;

    if (node == NUMA_NO_NODE)
        node = tsk_fork_get_node(orig);
    tsk = alloc_task_struct_node(node);
    if (!tsk)
        return NULL;

    stack = alloc_thread_stack_node(tsk, node);
    if (!stack)
        goto free_tsk;

    err = arch_dup_task_struct(tsk, orig);
    if (err)
        goto free_stack;

    tsk->stack = stack;
#ifdef CONFIG_SECCOMP
    /*
     * We must handle setting up seccomp filters once we're under
     * the sighand lock in case orig has changed between now and
     * then. Until then, filter must be NULL to avoid messing up
     * the usage counts on the error path calling free_task.
     */
    tsk->seccomp.filter = NULL;
#endif
}

C Tab Width: 8 Ln 342, Col 2 INS
```

```
#ifdef CONFIG_CC_STACKPROTECTOR
    tsk->stack_canary = get_random_int();
#endif

    /*
     * One for us, one for whoever does the "release_task()" (usually parent)
     */
    atomic_set(&tsk->usage, 2);
#ifdef CONFIG_BLK_DEV_IO_TRACE
    tsk->btrace_seq = 0;
#endif
    tsk->splice_pipe = NULL;
    tsk->task_frag.page = NULL;
    tsk->wake_q.next = NULL;

    account_kernel_stack(stack, 1);

    kcov_task_init(tsk);
    tsk->pf = 0;
    return tsk;

free_stack:
    free_thread_stack(stack);
free_tsk:
    free_task_struct(tsk);
    return NULL;
}
```

在 arch/x86/mm/fault.c 文件中定义变量 pfcount; 并修改 arch/x86/mm/fault.c 中 do_page_fault() 函数。每次产生缺页中断, do_page_fault() 函数会被调用, pfcount 变量

值递增 1, 记录系统产生缺页次数, current->pf 值递增 1 以记录当前进程产生缺页次数:

```
*/
enum x86_pf_error_code {
    PF_PROT      = 1 << 0,
    PF_WRITE     = 1 << 1,
    PF_USER      = 1 << 2,
    PF_RSVD      = 1 << 3,
    PF_INSTR     = 1 << 4,
    PF_PK        = 1 << 5,
};

/* modified by zijin */
unsigned long pfcount;

/*
 * Returns 0 if mmiotrace is disabled, or if the fault is not
 * handled by mmiotrace:
 */
static unsigned long address,
[
    struct vm_area_struct *vma;
    struct task_struct *tsk;
    struct mm_struct *mm;
    int fault, major = 0;
    unsigned int flags = FAULT_FLAG_ALLOW_RETRY | FAULT_FLAG_KILLABLE;

    /*modified by zijin*/
    pfcount++;
    current->pf++;

    tsk = current;
    mm = tsk->mm;

    /*
     * Detect and handle instructions that would cause a page fault for
     * both a tracked kernel page and a userspace page.
     */
    if (kmemcheck_active(regs))
```

最后, 实现一个自己的 System_call:

```
GNU nano 2.5.3      File: /linux-4.8/kernel/sys.c

#include <asm/io.h>
#include <asm/unistd.h>
extern unsigned long pfcount;
asmlinkage int sys_mysyscall(void)
{
    printk("#$#\n");
    printk("@all page fault: %lu\n", pfcount);
    printk("@current page fault: %lu\n", current->pf);
    struct task_struct *p = &init_task;
    for(;;(p=next_task(p))!=&init_task;)
    {
        printk("@The dirty page of process %d: %d\n", p->pid, p->nr_dirtied);
    }
    printk("$#$\n");
    return 0;
}

#ifdef SET_UNALIGN_CTL
#define SET_UNALIGN_CTL(a, b) (-EINVAL)

^G Get Help  ^O Write Out  ^W Where Is   ^K Cut Text   ^J Justify    ^C Cur Pos
^X Exit      ^R Read File  ^\ Replace    ^U Uncut Text ^T To Spell   ^_ Go To Line
```

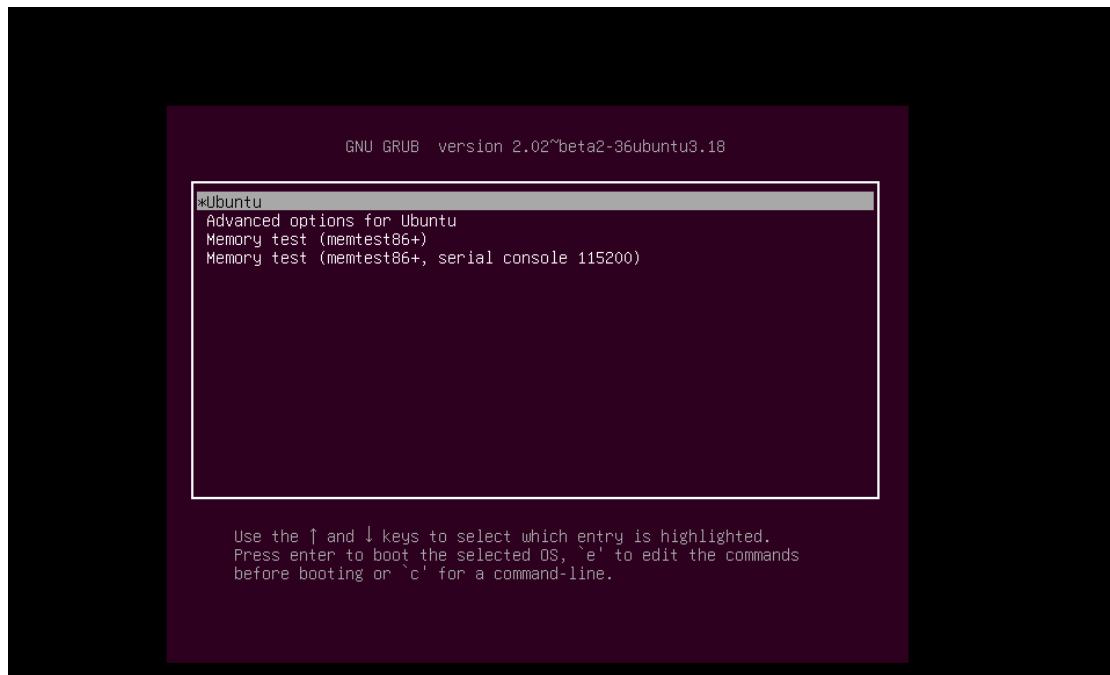
保存后执行 make。这个过程可能需要很长时间, 故我使用 make -j 以释放所有 CPU 资源:

```
danielshen@ubuntu: ~/linux-4.8
IHEX2FW firmware/emi62/midi.fw
IHEX2FW firmware/emi62/spdif.fw
IHEX firmware/kaweth/new_code.bin
IHEX firmware/kaweth/trigger_code.bin
IHEX firmware/kaweth/new_code_fix.bin
IHEX firmware/kaweth/trigger_code_fix.bin
IHEX firmware/ti_3410.fw
IHEX firmware/ti_5052.fw
IHEX firmware/mts_cdma.fw
IHEX firmware/mts_gsm.fw
IHEX firmware/mts_edge.fw
H16TOFW firmware/edgeport/boot.fw
H16TOFW firmware/edgeport/boot2.fw
H16TOFW firmware/edgeport/down2.fw
IHEX firmware/edgeport/down3.bin
H16TOFW firmware/edgeport/down.fw
IHEX2FW firmware/whiteheat_loader.fw
IHEX2FW firmware/keys span_pda/keys span_pda.fw
IHEX2FW firmware/whiteheat.fw
IHEX2FW firmware/keys span_pda/xircom_pgs.fw
IHEX firmware/cpia2/stv0672_vp4.bin
IHEX firmware/yam/9600.bin
IHEX firmware/yam/1200.bin
danielshen@ubuntu:~/linux-4.8$
```

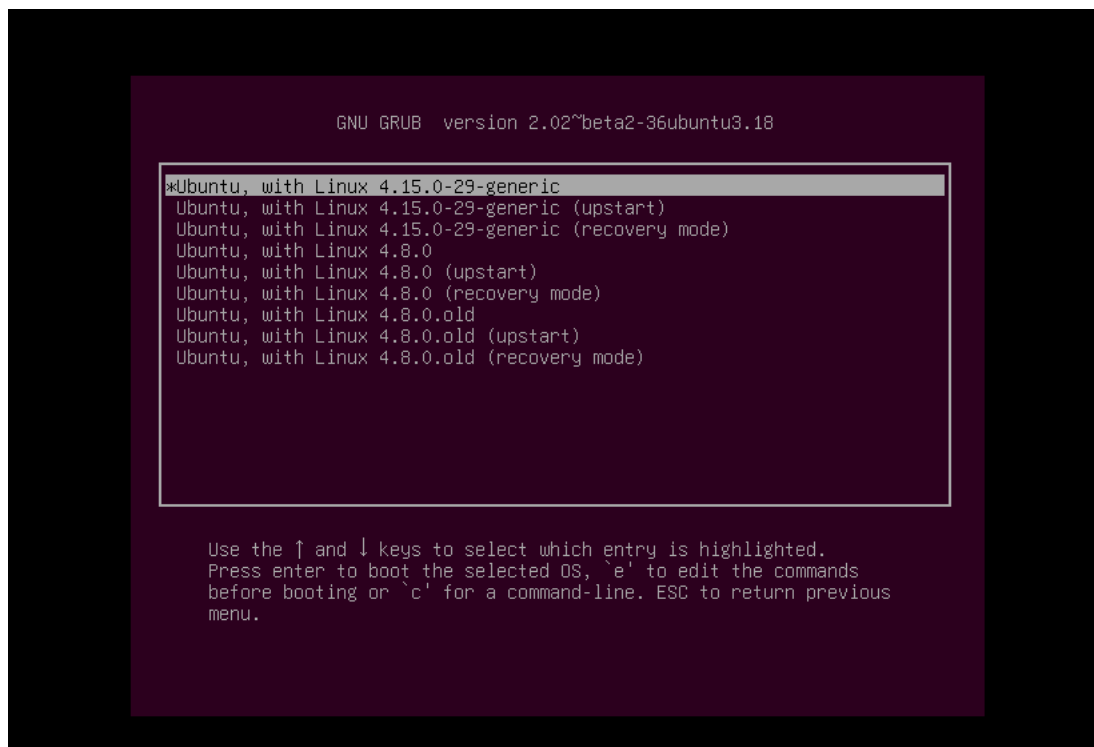
随后执行 make modules、make modules install 和 make install 命令进行最后配置:

```
danielshen@ubuntu: ~/linux-4.8
AS [M] arch/x86/crypto/sha512-mb/sha512_mb_mgr_submit_avx2.o
AS [M] arch/x86/crypto/sha512-mb/sha512_x4_avx2.o
LD [M] arch/x86/crypto/sha512-mb/sha512-mb.o
CC [M] arch/x86/events/intel/rapl.o
LD [M] arch/x86/events/intel/intel-rapl-perf.o
CC [M] arch/x86/events/intel/cstate.o
LD [M] arch/x86/events/intel/intel-cstate.o
CC [M] arch/x86/kernel/msr.o
CC [M] arch/x86/kernel/cpuid.o
CC [M] arch/x86/kernel/cpu/mcheck/mce-inject.o
CC [M] arch/x86/kvm/../../../../virt/kvm/kvm_main.o
CC [M] arch/x86/kvm/../../../../virt/kvm/coalesced_mmio.o
CC [M] arch/x86/kvm/../../../../virt/kvm/eventfd.o
CC [M] arch/x86/kvm/../../../../virt/kvm/irqchip.o
CC [M] arch/x86/kvm/../../../../virt/kvm/vfio.o
CC [M] arch/x86/kvm/../../../../virt/kvm/async_pf.o
CC [M] arch/x86/kvm/x86.o
^[[^A CC [M] arch/x86/kvm/mmu.o
CC [M] arch/x86/kvm/emulate.o
CC [M] arch/x86/kvm/i8259.o
CC [M] arch/x86/kvm/irq.o
CC [M] arch/x86/kvm/lapic.o
CC [M] arch/x86/kvm/i8254.o
```


配置完毕, 输入 sudo reboot 命令, 回车按下的一瞬间按住 esc 键, 直至进入内核切换界面



在这个界面选择 advanced options for ubuntu:



选择内核 Linux 4.8.0（无括号和.old 标注），回车启动，等待加载出图形界面：



3. 测试程序运行结果截图

编写读取结果的用户态程序。注意，本次使用的用户态程序和上一次实验使用的用户态程序是同一个：

```

#include <stdio.h>
#include <string.h>
#include <linux/unistd.h>
#include <sys/syscall.h>
#define __NR_mysyscall 223

int main(){
    int ch;
    int output=0;
    FILE *fp;
    /*
    fp = fopen("/var/log/kern.log","w");
    if(fp==NULL)
    { //若打不开则输出错误信息
        printf("No Permission!\n");
        return 0;
    }
    fclose(fp);
    */
    //打开日志文件
    syscall(__NR_mysyscall);
    fp=fopen("/var/log/kern.log","r");
    if(fp==NULL)
    { //若打不开则输出错误信息
        printf("No Permission!\n");
        return 0;
    }

    //读取日志文件
    fseek(fp,0,SEEK_SET);
    //文件指针重定位到文件头
    ch=fgetc(fp);
    //找到内核模块输出记录开头

```

编译运行该程序，结果（部分）如下：

HEAD	The dirty page of process 961: 0	The dirty page of process 1720: 0
all page fault: 547900	The dirty page of process 963: 0	The dirty page of process 1740: 0
current page fault: 411	The dirty page of process 964: 0	The dirty page of process 1763: 2
The dirty page of process 1: 0	The dirty page of process 965: 0	The dirty page of process 1779: 7
The dirty page of process 2: 0	The dirty page of process 966: 0	The dirty page of process 1810: 0
The dirty page of process 3: 0	The dirty page of process 967: 0	The dirty page of process 1825: 0
The dirty page of process 4: 0	The dirty page of process 968: 0	The dirty page of process 1826: 0
The dirty page of process 5: 0	The dirty page of process 969: 0	The dirty page of process 1828: 5
The dirty page of process 6: 0	The dirty page of process 970: 0	The dirty page of process 1829: 0
The dirty page of process 7: 0	The dirty page of process 971: 0	The dirty page of process 1833: 0
The dirty page of process 8: 0	The dirty page of process 972: 0	The dirty page of process 1835: 0
The dirty page of process 9: 0	The dirty page of process 973: 0	The dirty page of process 1854: 0
The dirty page of process 10: 0	The dirty page of process 989: 0	The dirty page of process 1856: 0
The dirty page of process 11: 0	The dirty page of process 990: 0	The dirty page of process 1871: 0
The dirty page of process 12: 0	The dirty page of process 998: 1	The dirty page of process 1883: 0
The dirty page of process 13: 0	The dirty page of process 1004: 0	The dirty page of process 1888: 0
The dirty page of process 14: 0	The dirty page of process 1006: 88	The dirty page of process 1893: 0
The dirty page of process 15: 0	The dirty page of process 1025: 11	The dirty page of process 1899: 0
The dirty page of process 16: 0	The dirty page of process 1027: 1	The dirty page of process 1914: 0
The dirty page of process 17: 0	The dirty page of process 1048: 0	The dirty page of process 1915: 0
The dirty page of process 18: 0	The dirty page of process 1203: 0	The dirty page of process 1926: 7
The dirty page of process 19: 0	The dirty page of process 1204: 0	The dirty page of process 1944: 0
The dirty page of process 20: 0	The dirty page of process 1209: 0	The dirty page of process 2007: 0
The dirty page of process 21: 0	The dirty page of process 1316: 0	The dirty page of process 2033: 0
The dirty page of process 22: 0	The dirty page of process 1330: 0	The dirty page of process 2039: 0
The dirty page of process 23: 0	The dirty page of process 1331: 0	The dirty page of process 2049: 0
The dirty page of process 24: 0	The dirty page of process 1348: 0	The dirty page of process 2056: 0
The dirty page of process 25: 0	The dirty page of process 1360: 1	The dirty page of process 2060: 9
The dirty page of process 26: 0	The dirty page of process 1380: 1	The dirty page of process 2067: 0
The dirty page of process 27: 0	The dirty page of process 1388: 0	The dirty page of process 2086: 0
The dirty page of process 28: 0	The dirty page of process 1389: 0	The dirty page of process 2109: 0
The dirty page of process 29: 0	The dirty page of process 1394: 1	The dirty page of process 2118: 0
The dirty page of process 30: 0	The dirty page of process 1396: 12	The dirty page of process 2151: 0
The dirty page of process 31: 0	The dirty page of process 1471: 0	
The dirty page of process 32: 0	The dirty page of process 1482: 0	
The dirty page of process 33: 0	The dirty page of process 1494: 0	
The dirty page of process 34: 0	The dirty page of process 1536: 1	
The dirty page of process 35: 0	The dirty page of process 1549: 0	
The dirty page of process 36: 0	The dirty page of process 1553: 0	

4. 结果分析

结合上述截图分析，大部分进程的脏页数都为0，少部分进程拥有0-100的脏页数，极少部分进程拥有略多于100的脏页数。数据基本合理，满足客观事实和实验要求，可以认为实验取得了成功。下面回答实验刚开始提出的问题：

- 多次运行 test 程序，每次运行 test 后记录下系统缺页次数和当前进程缺页次数，给出这些数据。test 程序打印的缺页次数是否就是操作系统原理上的缺页次数？有什么区别？

答：有一定区别。Test 程序中的“缺页次数”指的是调用 `do_page_fault` 的次数，真实的缺页次数应当不大于程序中的“缺页次数”，因为一次缺页可能调用多次 `do_page_fault` 函数，也可能有其他进程调用这一函数。

- 除了通过修改内核来添加一个系统调用外，还有其他的添加或修改一个系统调用的方法吗？如果有，请论述。

答：还可以通过内核模块的方式添加简单的系统调用。其原因是：编译内核的方式费时间，一般的 PC 机都要两三个小时，而且不方便调试，一旦出现问题前面的工作都前功尽弃。其基本步骤是：1. 找系统调用表在内存中的位置；2. 找一个空闲的系统调用号；3. 修改寄存器写保护位；4. 实现系统调用函数 5. 执行 `make` 后将编译好的模块插入内核。

- 对于一个操作系统而言，你认为修改系统调用的方法安全吗？请发表你的观点。

答：无论是采用插入模块的方式还是修改代码的方式编辑系统调用，严格来说都不是绝对安全的，甚至存在严重风险。以我自己的实验经历为例，在实验一中，如果内核模块代码的输出存在格式上的问题，可能破坏系统日志的结构；在实验二中，修改内核代码的行为更是危险重重，除了编译时间漫长，无法应对可能存在的系统宕机情况外，一个小小的 `runtime error` 就可能导致黑屏、卡死、GUI 无法加载等致命错误。在添加系统调用号的操作中，我为了使用 223 号调用而删除了一个已经存在的调用（获取系统时间的调用），这种修改如果不能得到专业人士的把关，势必会影响一些依赖此调用运行的应用程序，甚至操作系统模块的功能。

5. 源程序

系统调用函数：

```
1. extern unsigned long pfcount;
2. asm linkage int sys_mysyscall(void)
3. {
4.     printk("#$$\n");
5.     printk("@all page fault: %lu\n", pfcount);
6.     printk("@current page fault: %lu\n", current->pf);
7.     struct task_struct *p = &init_task;
8.     while((p=next_task(p))!=&init_task)
9.     {
10.         printk("@The dirty page of process %d: %d\n", p->pid, p->nr_dirtied);
11.     }
12.     rintk("#$$\n");
13.     return 0;
14. }
```

用户态程序（和实验一使用的是同一个）：

```
1. #include <stdio.h>
2. #include <string.h>
```



```
3.  #include <linux/unistd.h>
4.  #include <sys/syscall.h>
5.  #define __NR_mysyscall 223
6.
7.
8.  int main(){
9.      int ch;
10.     int output=0;
11.     FILE *fp;
12.     //打开日志文件
13.     syscall(__NR_mysyscall);
14.     fp=fopen("/var/log/kern.log", "r");
15.     if(fp==NULL)
16.         { //若打不开则输出错误信息
17.             printf("No Permission!\n");
18.             return 0;
19.         }
20.
21.     //读取日志文件
22.     fseek(fp,0,SEEK_SET);
23.     //文件指针重定位到文件头
24.     unsigned long offset = 0;
25.     ch=fgetc(fp);
26.     //找到内核模块输出记录开头
27.     FILE *begin = NULL;
28.     while(ch!=EOF)
29.     {
30.         while(ch!=EOF)
31.         {
32.             //输出记录开头的特殊标记为"##$"
33.             if(ch=='#')
34.             {
35.                 ch=fgetc(fp);
36.                 if(ch=='$')
37.                 {
38.                     ch=fgetc(fp);
39.                     if(ch=='#')
40.                     {
41.
42.                         printf("HEAD\n");
43.                         break;
44.                     }
45.                 }
46.             }
```

```

47.         ch=fgetc(fp);
48.     }
49.     //打印出内核模块的输出记录, @是用于识别每一行的特殊标记
50.     while(ch!=EOF)
51.     {
52.         if(ch == '@')
53.         {
54.             output = 1;
55.         }
56.         else if(ch == '\n')
57.         {
58.             printf("\n");
59.             output = 0;
60.         }
61.         else if(ch == '$')
62.             { //输出记录结尾的特殊标记为"$$"
63.             ch=fgetc(fp);
64.             if(ch=='#')
65.             {
66.                 ch=fgetc(fp);
67.                 if(ch=='$')
68.                 {
69.                     printf("END\n");
70.                     ch=fgetc(fp);
71.                     break;
72.                 }
73.             }
74.             if(output == 1 && ch != '@')
75.             {
76.                 printf("%c",ch);
77.             }
78.             ch=fgetc(fp);
79.         }
80.     }
81.     //关闭日志文件
82.     fclose(fp);
83.     return 0;
84. }

```

三、讨论、心得（20 分）

本次实验的操作步骤较为清晰,很多提示也都在实验指导中给出,看似要比实验一简单。但在完成该实验的过程中,我依然遇到了各种没有预料到的困难,有的至今没有找到最优的解决方案:

1. Ubuntu 的版本问题。本次实验要求的 Ubuntu 版本是 16.04，而我常用的版本是 18.x。
天真的我认为发行版的更新不会影响实验的推进，但残酷的现实给我上了一课——切换内核后，虚拟机无论如何也无法启动了。最终我选择重新安装 16.04 版本的 Ubuntu，将原来的 18.x 在移动硬盘中备份后从本机中删除了。整个过程耗费了我将近 3 个小时的时间。
2. 即使是正确的版本，问题依然存在。我使用新安装的 16.04 LTS 和在阿里云镜像下载的内核和补丁严格按照步骤推进实验，幻想这次可以一遍成功。但残酷的现实又给我上了一课——切换内核后，虚拟机启动了，但提示：

```
Gave up waiting for root device. Common problems:

- boot args (cat /proc/cmdline)

    - check root delay=(did the system wait long enough?)

    - check root=(did the system wait for the right device?)

- miss modules (cat /proc/modules; ls /dev)

ALERT! /dev/disk/by-uuid/acc3414d-926c-453c-b458-cf47088d77d2 does not exist. dropping to a shell!
```

- 期待中的图形界面并没有显示。检查所有文件修改操作无误后，上网查阅资料得知，可能是 kernel 中的某个 uuid 字段出了问题，尝试了一些方法，无果，遂怀疑是 vmware 配置的问题，但拷贝室友的虚拟机文件在本机的 vmware 上运行，并没有问题……最终我再一次重装了 Ubuntu，这次使用的全部是官网资源，然后成功启动了……所以是阿里云的问题？并不是，因为室友就是从阿里云上下载的内核，一点问题也没有……可能是之前打补丁或者输某些命令时出了问题吧。
3. 图形界面算是打开了，但还是有问题——脏页数太多了，有的进程有上万的脏页数，这很明显是不可能的。查阅资料得知，是 `tsk->pf=0` 的位置错了……这是因为 `arch_dup_task_struct(tsk, orig)` 函数的关系，如果看一下这个函数怎么实现的话就会知道：它直接把 `orig` 赋给了 `tsk`，所以如果 `pf` 在上面初始化，子进程的 `pf` 就还是父进程的，不满足要求。
 4. 还有一个坑是在 1 之前踩的：reboot 之后忘了长按 `esc` 切换内核……

总之这个实验可以说是处处皆坑。不管坑是老师故意让我们踩的还是确实没有说清楚，这次实验的收获还是蛮大的，可以说是“痛并快乐着”。以及，这是我上大学之后写的最长的实验心得……因为真的很有“心得”。